

A decorative graphic on the left side of the slide consisting of white lines and circles on a blue gradient background, resembling a circuit board or network diagram.

DISTRIBUTED AND CLOUD COMPUTING

LAB2 INTRODUCTION TO MPI COMMUNICATION MODELS

TROUBLESHOOTING

- Cannot open terminal (gnome-terminal) after installing Ubuntu in VirtualBox: possibly a problem in locale.

- Press **Ctrl+Alt+F3**, and log in
- Use admin permission to run the following commands:

```
locale-gen  
localectl set-locale LANG="en_US.UTF-8"  
reboot
```

- Wait for the reboot to complete
- Don't simply copy and paste the code! Type by yourself.
 - These two symbols are different: - -
 - E.g., when you run `mpirun -np 4 ./mpihw`

```
#include <mpi.h>
#include <stdio.h>
```

The hello-world example
No communication between MPI processes

```
int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);
    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);
    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
processor_name, world_rank, world_size);
    // Finalize the MPI environment.
    MPI_Finalize();
}
```

https://github.com/wesleykendall/mpitutorial/blob/gh-pages/tutorials/mpi-hello-world/code/mpi_hello_world.c

COMPILING AND RUNNING

- Use `mpicc` to compile an MPI application
 - Usage is the same as gcc: `mpicc foo.c -o foo`
- Use `mpirun` to launch the MPI application
 - `mpirun -np <num_process> <program>`
 - If you do not use `mpirun`, sub-processes can't be spawned.

MPI BASIC INTERFACE

```
int MPI_Init(int *argc, char ***argv)
```

- The first function call of any MPI application
- Initialize MPI environment
- Start of parallel part of the application

MPI BASIC INTERFACE


- `int MPI_Comm_rank(MPI_Comm comm, int *rank)`
 - Get the rank of the current process in a communicator.
- `int MPI_Comm_size(MPI_Comm comm, int *size)`
 - Get the size (num. of processes) in a communicator.
- `MPI_COMM_WORLD`
 - The default communicator, groups all processes at start. There's also a communicator `MPI_COMM_SELF` provided by MPI, to communicate with the process itself.

MPI BASIC INTERFACE

- `int MPI_Finalize()`
 - The last function call of MPI application.
 - Terminate the MPI application, also the sub-processes created by mpirun
 - End of parallel part of MPI application. Serial code can still be run in the main process.

AN EXAMPLE: PING-PONG

- The hello-world example does not involve communication between processes
- Here we consider an example that allows two MPI processes to play ping-pong
- MPI processes send messages to each other




```
const int PING_PONG_LIMIT = 10;
```

```
.....
```

```
int ping_pong_count = 0;
```

```
int partner_rank = (world_rank + 1) % 2;
```

```
while (ping_pong_count < PING_PONG_LIMIT) {
```

```
    if (world_rank == ping_pong_count % 2) {
```

```
        // Increment the ping pong count before you send it
```

```
        ping_pong_count++;
```

```
        MPI_Send(&ping_pong_count, 1, MPI_INT, partner_rank, 0, MPI_COMM_WORLD);
```

```
        printf("%d sent and incremented ping_pong_count %d to %d\n",
```

```
                world_rank, ping_pong_count, partner_rank);
```

```
    } else {
```

```
        MPI_Recv(&ping_pong_count, 1, MPI_INT, partner_rank, 0,
```

```
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```

```
        printf("%d received ping_pong_count %d from %d\n",
```

```
                world_rank, ping_pong_count, partner_rank);
```

```
    }
```

```
}
```

Example program: ping-pong

It implements message passing between two MPI processes

```
const int PING_PONG_LIMIT = 10;
```

```
.....
```

```
int ping_pong_count = 0;
```

```
int partner_rank = (world_rank + 1) % 2;
```

```
while (ping_pong_count < PING_PONG_LIMIT) {
```

```
    if (world_rank == ping_pong_count % 2) {
```

```
        // Increment the ping pong count before you send it
```

```
        ping_pong_count++;
```

```
        MPI_Send(&ping_pong_count, 1, MPI_INT, partner_rank, 0, MPI_COMM_WORLD);
```

```
        printf("%d sent and incremented ping_pong_count %d to %d\n",
```

```
               world_rank, ping_pong_count, partner_rank);
```

```
    } else {
```

```
        MPI_Recv(&ping_pong_count, 1, MPI_INT, partner_rank, 0,
```

```
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```

```
        printf("%d received ping_pong_count %d from %d\n",
```

```
               world_rank, ping_pong_count, partner_rank);
```

```
    }
```

```
}
```

Example program: ping-pong

It implements message passing between two MPI processes

world_rank	partner_rank
1	0
0	1

Sender buffer

Num. of elements
to be sent


Data type

Rank of
destination process

Tag

Communicator

Two MPI processes playing ping-pong



```
~ mpirun -np 2 ./mpi-ping-pong
0 sent and incremented ping_pong_count 1 to 1
1 received ping_pong_count 1 from 0
1 sent and incremented ping_pong_count 2 to 0
0 received ping_pong_count 2 from 1
0 sent and incremented ping_pong_count 3 to 1
0 received ping_pong_count 4 from 1
0 sent and incremented ping_pong_count 5 to 1
0 received ping_pong_count 6 from 1
0 sent and incremented ping_pong_count 7 to 1
0 received ping_pong_count 8 from 1
0 sent and incremented ping_pong_count 9 to 1
0 received ping_pong_count 10 from 1
1 received ping_pong_count 3 from 0
1 sent and incremented ping_pong_count 4 to 0
1 received ping_pong_count 5 from 0
1 sent and incremented ping_pong_count 6 to 0
1 received ping_pong_count 7 from 0
1 sent and incremented ping_pong_count 8 to 0
1 received ping_pong_count 9 from 0
1 sent and incremented ping_pong_count 10 to 0
```

MPI SEND/RECEIVE

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)
```

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)
```

/**OUT buf	Sender buffer
*IN count	Num. of elements to be sent
*IN datatype	Data type
*IN source/dest	Rank of source/destination process
*IN tag Message	Tag
*IN comm	Communicator
*OUT status	Contains the actual received message */

MPI DATATYPE

- MPI supports various data types to be send among processes.
- In complex MPI applications, we typically use MPI_BYTE to communicate with custom protocols.

MPI datatype	C datatype
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED_INT	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	

MPI COMMUNICATION MODE

- **Standard Mode, Buffered Mode, Synchronous Mode, Ready Mode**
- They have the same set of parameters
- Differences: The method of sending message and the state of receiver
- Locality of mode: If the mode requires communicating with other processes.
 - Local: Completion of procedure **depends only on local process**
 - Non-local: Completion of procedure needs to execute some MPI procedure **on another process**

MPI COMMUNICATION MODE

- Standard mode
 - In standard mode, **MPI library itself determines if the data will be buffered.**
 - Buffered: Copy the data into a buffer and return immediately. The sending will be done by MPI in background.
 - Non-buffered: Return when the data has completed sending.
 - Standard mode is **non-local**, since it may need to communicate with recipient to complete sending.

MPI COMMUNICATION MODE

- Buffered mode
 - In buffered mode, **MPI copies the data to a provided buffer and return immediately**. The sending is done by MPI in background.
 - Buffered mode is local.

MPI COMMUNICATION MODE

- Synchronous mode
 - Synchronous mode only returns when the recipient has started receiving message.
 - Synchronous mode is non-local.

MPI COMMUNICATION MODE

- Ready mode
 - Ready mode ensures the recipient is at ready state (waiting for receiving message), or it will raise an error
 - Ready mode is non-local.

MPI BLOCKING AND NON-BLOCKING COMMUNICATION

- Blocking communication: The function waits for operation to complete, or at least is securely copied by MPI library.
- Non-blocking communication: Always returns immediately, **the actual operation is completed by MPI in background**. User must ensure the operation is completed before doing the next, or there may be conflicts.

SEND/RECEIVE FUNCTIONS

SEND	Blocking	Nonblocking
Standard	<code>mpi_send</code>	<code>mpi_isend</code>
Ready	<code>mpi_rsend</code>	<code>mpi_irsend</code>
Synchronous	<code>mpi_ssend</code>	<code>mpi_issend</code>
Buffered	<code>mpi_bsend</code>	<code>mpi_ibsend</code>
RECEIVE	Blocking	Nonblocking
Standard	<code>mpi_recv</code>	<code>mpi_irecv</code>

COLLECTIVE COMMUNICATION


- Collective communication refers to the communication among multiple (3 or more) processes.
- One to many(1-N), many to one(N-1), many to many(N-N)
- In 1-N, N-1, the single process is called *root*.
- Functionalities: Communicate, synchronization, computation

COLLECTIVE COMMUNICATION


- `int MPI_Barrier(MPI_Comm comm)`
 - Synchronization
- `int MPI_Bcast(void* buf, int count, MPI_Datatype datatype, int root, MPI_Comm comm)`
 - Broadcast message to all processes (including self)
- `int MPI_Gather(void * sendbuf, int sendcount, MPI_Datatype sendtype, void * recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)`
- `int MPI_Scatter(void * sendbuf, int sendcount, MPI_Datatype sendtype, void * recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)`

1-N AND N-1 OPERATIONS


MPI_Bcast



MPI_Scatter







MPI_Gather



COLLECTIVE COMMUNICATION


- `int MPI_Reduce(void * sendbuf, void * recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)`
- `int MPI_Allreduce(const void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)`
- `int MPI_Allgather(void * sendbuf, int sendcount, MPI_Datatype sendtype, void * recvbuf, int recvcount, MPI_Datatype recvtype, MPI_Comm comm)`
- `int MPI_Alltoall(void * sendbuf, int sendcount, MPI_Datatype sendtype, void* recvbuf, int recvcount, MPI_Datatype recvtype, MPI_Comm comm)`

N-N OPERATIONS



RUNNING MPI PROGRAMS ON A VM CLUSTER

- MPI programs run on clusters that are composed of multiple machines
- We can emulate such a cluster with multiple VMs
- Our task: creating a cluster of VMs (here we use two VMs), and run an MPI program on this cluster
- Steps:
 - Set up the network connection between the two VMs
 - Enable password-less SSH connection between the two VMs
 - Compile the MPI program and deploy it on each VM
 - Run the MPI program

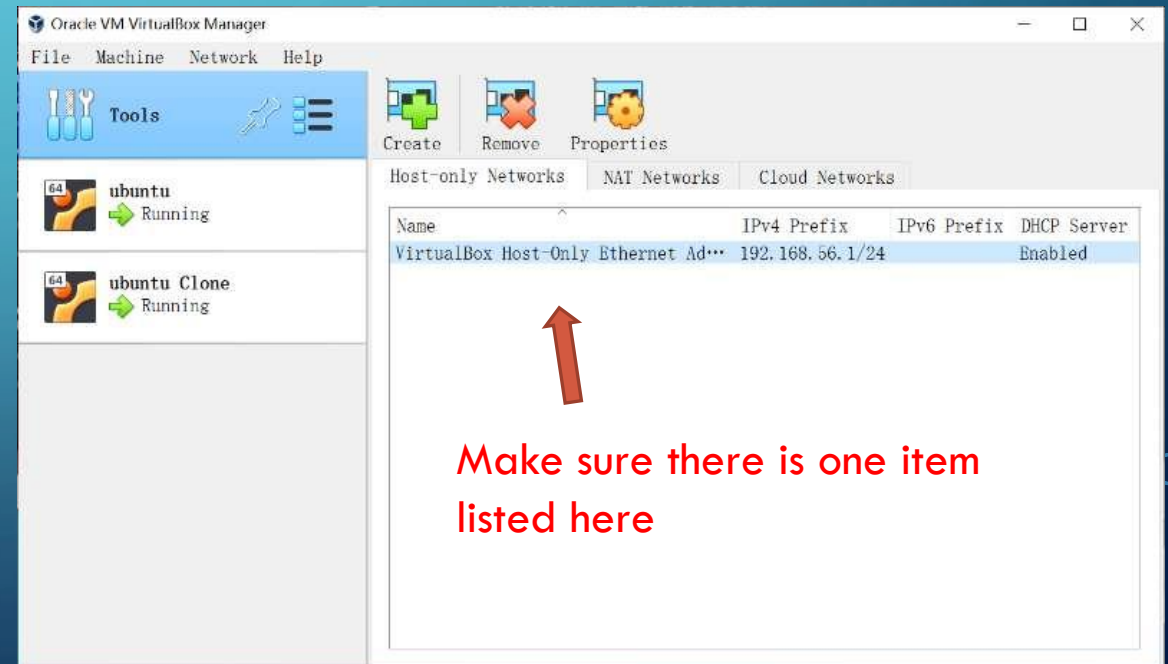
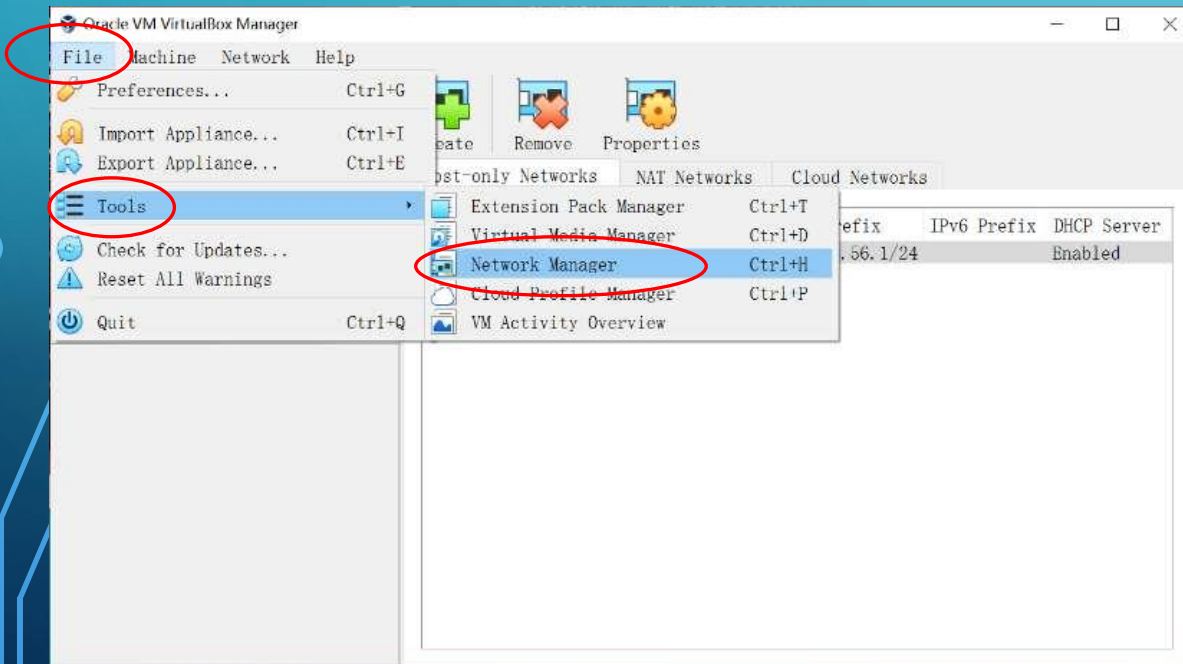


RUNNING MPI PROGRAMS ON A VM CLUSTER

- Prepare two VMs. Each VM should have the followings installed:
 - SSH server
 - `sudo apt install openssh-server`
 - ifconfig
 - `sudo apt install net-tools`

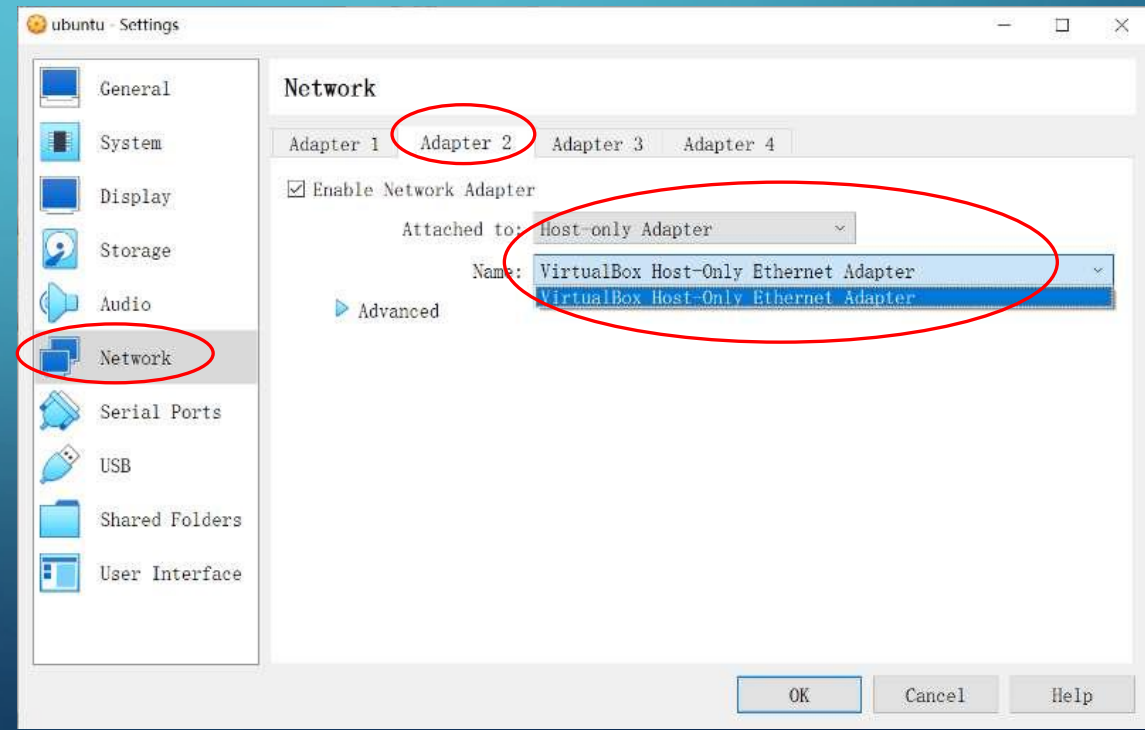

RUNNING MPI PROGRAMS ON A VM CLUSTER

- Make sure there is an existing host-only network adapter
 - If not, create one.



RUNNING MPI PROGRAMS ON A VM CLUSTER

- Shut down you VMs
- Set **EACH** of your VM to that host-only adapter (see the previous slide)



RUNNING MPI PROGRAMS ON A VM CLUSTER

- On each VM
 - Start the SSH server:
 - `sudo systemctl restart sshd.service`
 - Generate SSH keys
 - `ssh-keygen -t rsa`
 - Use `ifconfig` to check its own IP address
 - Copy the SSH key to the other VM
 - `ssh-copy-id username@remotehostIP`

RUNNING MPI PROGRAMS ON A VM CLUSTER

- The figures show how to enable password-less SSH from VM2 to VM1.
- You should also set password-less SSH from VM1 to VM2.

On VM1

```
nan@ubuntu:~$ ifconfig
enp0s3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.2.15 netmask 255.255.255.0 broadcast 10.0.2.255
    inet6 fe80::d625:cfc:c6ad:b91b prefixlen 64 scopeid 0x20<link>
    ether 08:00:27:02:b0:59 txqueuelen 1000 (Ethernet)
    RX packets 2778 bytes 3813532 (3.8 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1427 bytes 106400 (106.4 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

enp0s8: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.56.101 netmask 255.255.255.0 broadcast 192.168.56.255
    inet6 fe80::242a:9431:19fd:ed47 prefixlen 64 scopeid 0x20<link>
    ether 08:00:27:e0:71:3d txqueuelen 1000 (Ethernet)
    RX packets 197 bytes 37397 (37.3 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 195 bytes 33715 (33.7 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 216 bytes 23876 (23.8 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 216 bytes 23876 (23.8 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

On VM2

```
nan@ubuntu:~$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/nan/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/nan/.ssh/id_rsa
Your public key has been saved in /home/nan/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:0c5RN8H4z+8eyFfiawhHnOCX5A3Bm/BvVy71ShqLRkg nan@ubuntu
The key's randomart image is:
+---[RSA 3072]-----+
|      .+..          |
|      .o=O*         |
|                      |
|                      |
|                      |
+---+-----+-----+

```

```
nan@ubuntu:~$ ssh-copy-id nan@192.168.56.101
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter o
ut any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are prompte
d now it is to install the new keys
nan@192.168.56.101's password:
```

Number of key(s) added: 1

Now try logging into the machine, with: "ssh 'nan@192.168.56.101'"
and check to make sure that only the key(s) you wanted were added.

```
nan@ubuntu:~$ ssh nan@192.168.56.101
Welcome to Ubuntu 22.04.1 LTS (GNU/Linux 5.19.0-32-generic x86_64)
```

```
* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:        https://ubuntu.com/advantage
```

267 updates can be applied immediately.
128 of these updates are standard security updates.
To see these additional updates run: apt list --upgradable

Last login: Tue Feb 21 01:38:51 2023 from 192.168.56.102

```
nan@ubuntu:~$ exit
```

logout

Connection to 192.168.56.101 closed.

```
nan@ubuntu:~$
```

SSH without password

RUNNING MPI PROGRAMS ON A VM CLUSTER

- Compile the MPI program, and copy it to the two VMs
- Run the program by specifying the hosts on either of the VM
 - `mpirun -np 2 --host <IP_of_VM1>,<IP_of_VM2> <your_program>`


<https://mpitutorial.com/tutorials/running-an-mpi-cluster-within-a-lan/>

How to install and configure OpenSSH


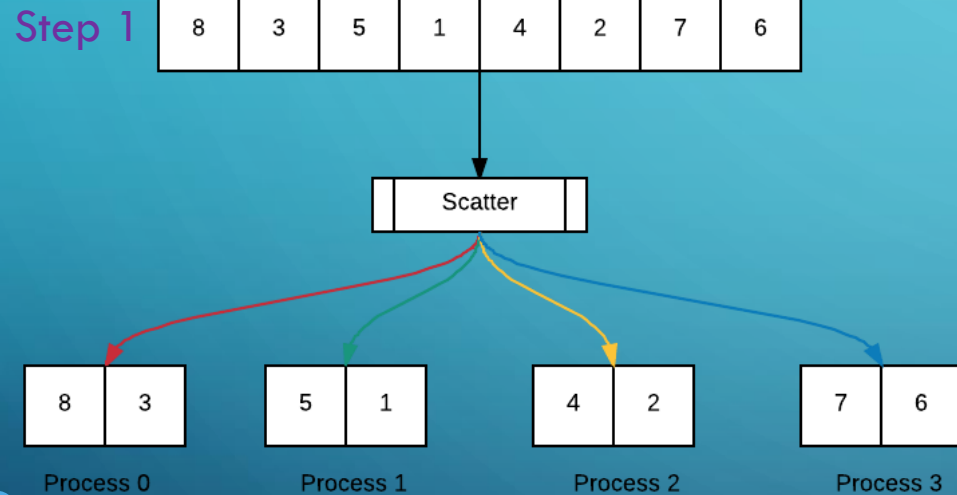
<https://ubuntu.com/server/docs/service-openssh>

PARALLELIZE COMPUTATION – MERGE SORT

- Merge sort is a classic divide-and-conquer sorting algorithm
- Process: Divide list into unsorted sub-lists, then sort sub-lists, finally merge all sorted sub-lists.
- Good candidate for parallelize: Sorting of sub-lists are independent!



PARALLELIZE COMPUTATION – MERGE SORT



TASKS TODAY

- Set up the VM cluster
- Run the ping-pong program on the VM cluster

PRACTICE: MATRIX MULTIPLICATION

- Rule for matrix multiplication $(AB)_{ij} = \sum_{r=1}^n a_{ir}b_{rj} = a_{i1}b_{1j} + a_{i2}b_{2j} + \cdots + a_{in}b_{nj}$
- Think about how to parallelize the computation process, and write an MPI application to do this.
- Can you parallelize the computation and also minimize the memory usage?
- Useful link: <https://mpitutorial.com/tutorials/>

