

Distributed Systems


Interprocess Communication



Network Architecture

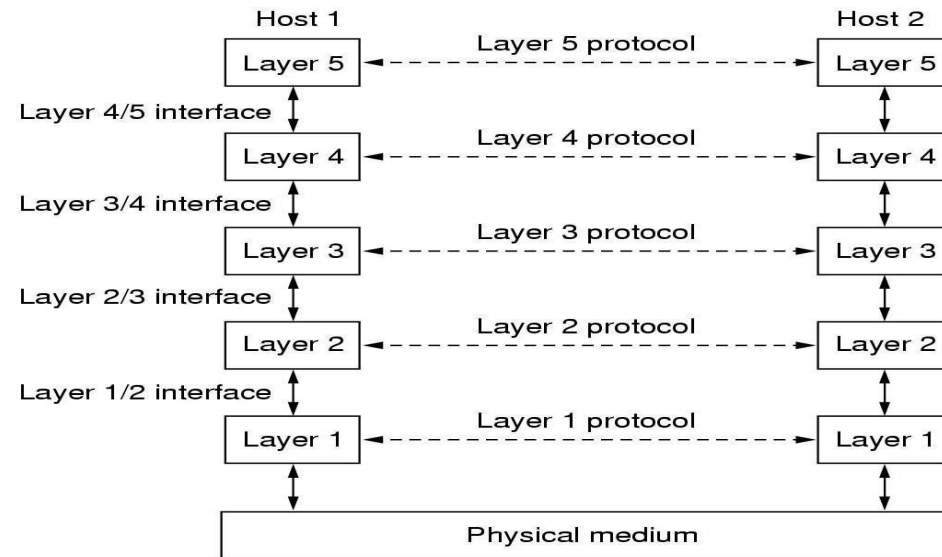
- Network Architecture: framework for designing and implementing Networks
- Components:
 - Software
 - Protocols
 - Services
 - Hardware
 - Transmission technology, media and devices
 - Scale: LANs, MANs, and WANs
 - Topology

Subnet



Layered Architecture

- The purpose of each layer is to offer a communication services to higher level layers
- Each layer has two interfaces
 - peer-to-peer interface
 - defines the form and types of messages exchanged between peers (indirect communication)
 - service interface
 - defines the primitives (operations) that a layer provides to the layer above it
- Layering is non-linear



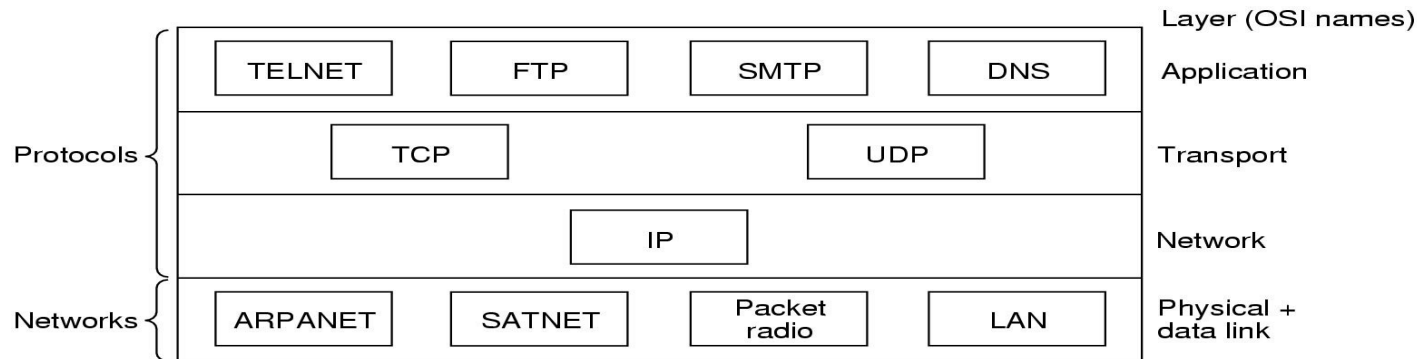
Protocols

- The functionality encapsulated within each layer is called *Protocol*
- The Protocol refers to both
 - the abstract peer-to-peer & service interfaces and
 - the objects that implement those interfaces
- Protocol vs. Service
 - Service is the set of primitives provided to the higher layer
 - Protocol defines the implementation of these primitives
- Protocol stack

ISO OSI Architecture

- International Standards Organisation (ISO)
- *Physical*: transmission of raw bits onto the communications medium
- *Data link*: reliable transmission of frames, flow control, arbitration
- *Network*: packet switching, routing congestion control
- *Transport*: process-to-process channel, node-to-node connection, provides user services, flow control, multiplexing
- *Session*
- *Presentation*
- *Application*

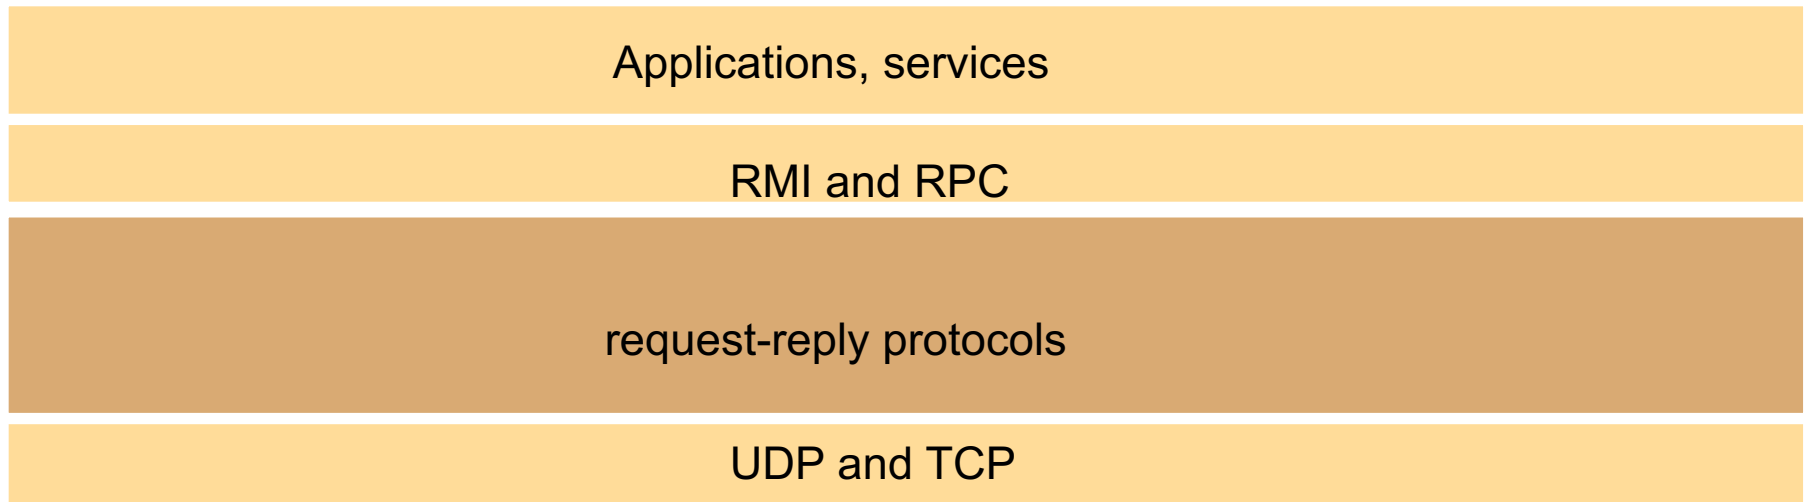
Internet Architecture (TCP/IP)



- Host-to-Network Layer (OSI Physical and Data link layers)
- Internet Layer (OSI Network layer - Internet Protocol/IP)
- **Transport Layer**
 - Transmission Control Protocol - TCP
 - User Datagram Protocol - UDP)
- Application Layer

Accessing Transport Layer

- How application programs use protocol software (invoke the services) of transport layer to communicate across networks




Inter-process Communication

- Message passing model: send - receive
- Synchronous communications : blocking
- Asynchronous communications
 - Blocking (receive) vs non blocking (send, receive)
- Reliability
 - Validity: messages are guaranteed to be delivered despite of packets being lost
 - Integrity: messages must arrive uncorrupted and without duplication
- Ordering: message send and delivery

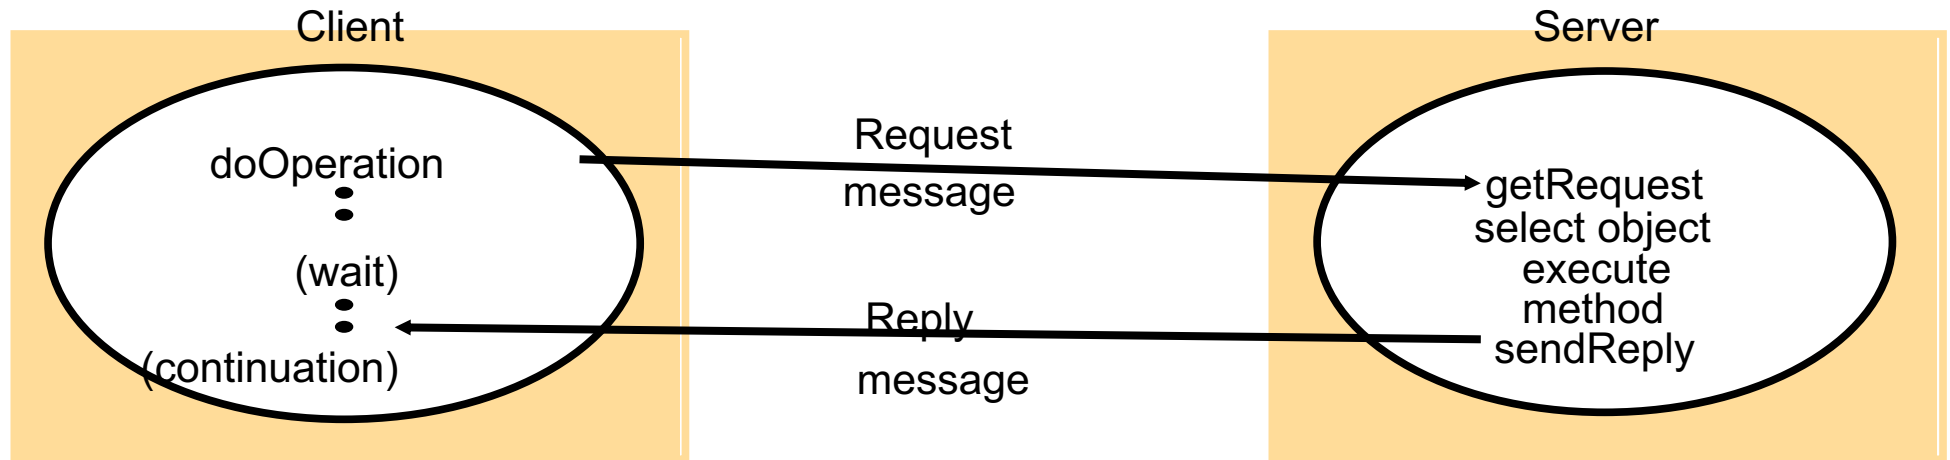
Sockets and ports

- Sockets provide an abstract endpoint for inter-process communication (UDP, TCP)
- Destination:
 - internet address, port, process id
- Port: message destination within a computer (integer) –Internet addresses: names rather than numbers (2^{16})
- A socket must be bound to a local port and an internet address
- A process may use multiple ports but cannot share ports (exception: multicast)



The Request-Reply Protocol

- Processes communicate using a protocol



```
public byte[] doOperation (RemoteObjectRef o, int methodId, byte[] arguments)
```

sends a request message to the remote object and returns the reply.

The arguments specify the remote object, the method to be invoked and the arguments of that method.

```
public byte[] getRequest ();
```

acquires a client request via the server port.

```
public void sendReply (byte[] reply, InetAddress clientHost, int clientPort);
```

sends the reply message reply to the client at its Internet address and port.

Issues

- Message identifiers (sequence+sender id)
- Failure Model
- Timeouts
- Duplicate request messages
- Lost reply messages
- History for retransmissions

Java API for Internet addresses

- Class *InetAddress*
 - uses DNS (Domain Name System)

```
InetAddress aComputer =  
    InetAddress.getByName("gromit.cs.bham.ac.uk");
```

- throws *UnknownHostException*
- encapsulates detail of IP address (4 bytes for IPv4 and 16 bytes for IPv6)

UDP Sockets

- Receive method returns the IP address+port of sender so that a reply can be sent
- Message size: receiving process specifies a byte-array in which to receive the message (IP: 2^{16} bytes, usually 8kB)
- Non-blocking send & blocking receive (timeouts and threads for deadlocks)
- Receive *from any*

UDP Sockets – Java API

- *DatagramPacket* class provides
 - a constructor that makes an instance of a message: [message byte array, length, IP address, port no]
 - A constructor for use when receiving a message
 - Methods: *getData*, *getPort*, *getAddress*
- *DatagramSocket* class provides
 - A constructor that takes a port as an argument for use by processes that need to use a port
 - A no-argument constructor that allows the system to choose a free local port
 - *SocketException* if problem with port
 - Methods:
 - *send*: argument an instance of *DatagramPacket*
 - *Receive*: argument an empty *DatagramPacket*
 - *setSoTimeout*: *InterruptedIOException* when expire
 - *connect*: connect a socket to a particular remote port and IP address, the socket is able to send to and receive from that address

Java API for Datagrams


- Simple send/receive, with messages possibly lost/out of order
- Class *DatagramPacket*

message (=array of bytes)	message length	Internet addr	port no
---------------------------	----------------	---------------	---------

- packets may be transmitted between sockets
- packets truncated if too long
- provides *getData*, *getPort*, *getAddress*

example...

- UDP Client
 - sends a message and gets a reply
- UDP Server
 - **repeatedly** receives a request and sends it back to the client



UDP client sends a message to the server and gets a reply

```
import java.net.*;
import java.io.*;
public class UDPClient{
    public static void main(String args[]){
        // args give message contents and server hostname
        DatagramSocket aSocket = null;
        try {
            aSocket = new DatagramSocket();
            byte [] m = args[0].getBytes();
            InetAddress aHost = InetAddress.getByName(args[1]);
            int serverPort = 6789;
            DatagramPacket request = new DatagramPacket(m, args[0].length(), aHost, serverPort);
            aSocket.send(request);
            byte[] buffer = new byte[1000];
            DatagramPacket reply = new DatagramPacket(buffer, buffer.length);
            aSocket.receive(reply);
            System.out.println("Reply: " + new String(reply.getData()));
        }catch (SocketException e){System.out.println("Socket: " + e.getMessage());}
        }catch (IOException e){System.out.println("IO: " + e.getMessage());}
    }finally {if(aSocket != null) aSocket.close();}
}
}
```

UDP server repeatedly receives a request and sends it back to the client

```
import java.net.*;
import java.io.*;
public class UDPServer{
    public static void main(String args[]){
        DatagramSocket aSocket = null;
        try{
            aSocket = new DatagramSocket(6789);
            byte[] buffer = new byte[1000];
            while(true){
                DatagramPacket request = new DatagramPacket(buffer, buffer.length);
                aSocket.receive(request);
                DatagramPacket reply = new DatagramPacket(request.getData(),
                    request.getLength(), request.getAddress(), request.getPort());
                aSocket.send(reply);
            }
        }catch (SocketException e){System.out.println("Socket: " + e.getMessage());}
        }catch (IOException e) {System.out.println("IO: " + e.getMessage());}
    }finally {if(aSocket != null) aSocket.close();}
}
```

TCP Sockets

- API provides an abstraction of a stream of bytes to which data may be written and from which data may be read. The abstraction hides the following:
 - Message sizes. The application can choose how much data it writes to a stream. TCP may fragment it to smaller packets
 - Lost messages
 - Flow control
 - Message duplication and ordering
 - Message destinations: once a connection is established, processes simply read from and write to a stream (no need for IP addresses and ports)
- Client-server model during connection:
 - client creates a stream socket bound to a port and asks for a connection to a server port
 - server creates a listening socket bound a a port and waits to *accept* connect requests
- During operation each socket is both for input and output
- Close a socket when no more data to write

TCP Sockets – Java API

- *ServerSocket* class: create a socket at a server port for listening for connect requests
 - Method *accept*: wait until there is a *connect* request in the queue, then create an instance of *Socket*
- *Socket* class: client uses a constructor which creates a socket (specifying the DNS hostname and port of a server) and connects it to the specified remote server (*UnknownHostException*, *IOException*)
 - Methods: *getInputStream* and *getOutputStream*
 - Return types are abstract classes that define methods for reading (*InputStream*) and writing (*OutputStream*) bytes
 - Return values can be used as the arguments of constructors for suitable input and output streams.

TCP client makes connection to server, sends request and receives reply

```
import java.net.*;
import java.io.*;
public class TCPClient {
    public static void main (String args[]) {
        // arguments supply message and hostname of destination
        Socket s = null;
        try{
            int serverPort = 7896;
            s = new Socket(args[1], serverPort);
            DataInputStream in = new DataInputStream( s.getInputStream());
            DataOutputStream out =
                new DataOutputStream( s.getOutputStream());
            out.writeUTF(args[0]);           // UTF is a string encoding see Sn 4.3
            String data = in.readUTF();
            System.out.println("Received: "+ data) ;
        }catch (UnknownHostException e){
            System.out.println("Sock:"+e.getMessage());
        }catch (EOFException e){System.out.println("EOF:"+e.getMessage());}
        }catch (IOException e){System.out.println("IO:"+e.getMessage());}
    }finally {if(s!=null) try {s.close();}catch (IOException e){System.out.println("close:"+e.getMessage());}}
    }
}
```

TCP server makes a connection for each client and then echoes the client's request

```
import java.net.*;
import java.io.*;
public class TCPServer {
    public static void main (String args[]) {
        try{
            int serverPort = 7896;
            ServerSocket listenSocket = new ServerSocket(serverPort);
            while(true) {
                Socket clientSocket = listenSocket.accept();
                Connection c = new Connection(clientSocket);
            }
        } catch(IOException e) {System.out.println("Listen :"+e.getMessage());}
    }
}
```

continued

```
class Connection extends Thread {
    DataInputStream in;
    DataOutputStream out;
    Socket clientSocket;
    public Connection (Socket aClientSocket) {
        try {
            clientSocket = aClientSocket;
            in = new DataInputStream( clientSocket.getInputStream());
            out =new DataOutputStream( clientSocket.getOutputStream());
            this.start();
        } catch(IOException e) {System.out.println("Connection:"+e.getMessage());}
    }
    public void run(){
        try {                                // an echo server
            String data = in.readUTF();
            out.writeUTF(data);
        } catch(EOFException e) {System.out.println("EOF:"+e.getMessage());}
        } catch(IOException e) {System.out.println("IO:"+e.getMessage());}
        } finally{ try {clientSocket.close();}catch (IOException e){/*close failed*/}}
    }
}
```


Group Communication:

- Multicast: an operation that sends a single message from one process to each of the members of a group of processes
- Fault tolerance based on replicated services
- Finding the discovery servers in spontaneous networking
- Better performance through replicated data
- Propagation of event notifications

IP multicast

- multicast group is specified by a class D Internet address
 - membership is dynamic, to join make a socket
 - programs using multicast use UDP and send datagrams to multicast addresses and (ordinary) port
 - (For example of Java code see book)